


Second Brain for Claude

LLM Wiki vs RAG

Florian Steiner

Claude Community Ambassador Munich

Monday, 21 September 2026 · 11:00 to 12:00 · Auditorium



My assistant searched for ten minutes and came back with nothing.

Everything it needed sat in a folder it did not know existed.

Not a memory problem. A map problem.

Context is finite

Every token spends attention.

Loading everything is not an option.

So give the agent a map

Know where to look.

Then load only what the question needs.

What RAG does well

Large corpora

Scales far past what fits in any context window.

Unknown questions

Retrieves at question time, including the long tail.

Fresh documents

Add a file to the index. No retraining.

Where RAG loses context

Chunk boundaries

"The company's revenue grew by 3%." Which company? Which quarter?

ANTHROPIC 2024

No canonical fact

Last year's price list ranks as high as this year's.

OWN OBSERVATION

Contradictions side by side

Both versions reach the prompt. The model must pick one.

WANG ET AL. 2025


Summaries on top of retrieval help

retrieval searches
all levels

summarise,
repeat

embed +
cluster

chunks



RAPTOR builds its tree with embeddings and clustering, then retrieves across summary levels.

+20%


absolute accuracy on QuALITY with GPT-4, over the previous best result

Evidence that summaries on top of retrieval help. Not evidence that hierarchy beats vector search.

Retrieve every time, or compile once


RAG

work at question time



LLM wiki


work at ingest time



“A persistent, compounding artifact.”

Karpathy, on the wiki

Root map: one small file, always loaded




The agent reads the map first, then descends only where the question leads.

Canonical fact sheet: one home per fact

FACTS.md			example
FACT	VALUE	SOURCE	AS OF
Support hours	9:00 to 18:00	ops-handbook.md	2026-05
Standard warranty	3 years	policy-2026.md	2026-03
Current price list	v4	pricing-v4.md	2026-08

Other pages link here instead of copying. Every value carries its source.

Ingest: read once, update many pages



10 to 15

wiki pages a single source might touch.

Karpathy's own research wiki: about 100 articles and 400K words, navigated through index files and summaries.

Lint: a health check for the wiki

☒ Contradictions between pages

☒ Stale claims a newer source superseded

☒ Duplicates: one thing, two pages

OWN ADDITION

☒ Orphans: pages nothing links to

OWN OBSERVATION

Skip lint, and a wrong fact gets
filed back, then cited again.

1 Ingest

2 Lint

3 Fact sheet

1

Ingest a new source

A new document for **Sotarena S.L.** lands in sources/inbox/. Watch which wiki pages change.

1 Ingest

2 Lint

3 Fact sheet

2

Lint finds a contradiction, a duplicate and a stale fact

Contradiction

Two values for one fact

Duplicate

One entity, two pages

Stale fact

Superseded by a newer source

1 Ingest

2 Lint

3 Fact sheet

3

The fact sheet shows the winning value, with its source

sotarena / FACTS.md			after lint
FACT	WINNING VALUE	SOURCE	REPLACES
[fact]	[value]	[file, date]	[old value]

Compile what matters. Retrieve the rest.

Root map

always loaded

LLM wiki

curated, recurring, canonical

RAG over raw sources

large, unforeseen, fresh



Promote recurring answers
into the wiki



Fall back to retrieval when no
page exists

Under about 200K tokens? Load it all. No retrieval needed.

TAKE IT HOME

The starter repo

`github.com/ProduktEntdecker/
enterprise-memory-starter`

Repository opens before the session

Questions?

11:50 to 12:00

Florian Steiner · Claude Community Ambassador Munich



Fallback: Ingest a new source

[SCREENSHOT]

Fallback: Lint finds a contradiction, a duplicate and a stale fact

[SCREENSHOT]

Fallback: The fact sheet shows the winning value and its source

[SCREENSHOT]